



uOttawa

## School of Information Technology and Engineering

CEG 4188: Higher Layer Network Protocols

### Assignment 1 (bonus question)

#### **Bonus question (5 marks): Programming assignment**

**(This assignment is taken from a course taught by Prof. Scott Shenker @ Berkley)**

For this question you will need to build a simple application that connects users over a network: a chat server. Your chat server will allow users to converse in different channels. Users can create and join channels; once a user is in a particular channel, all messages that he/she sends will be relayed to all other users in that channel.

This assignment should be implemented in Python 2. This assignment will introduce you to the socket programming API.

**Download and use the following help files from BBL:**

- `client_split_messages.py` is intended to help you ensure your server is correctly buffering messages, and is described in more detail below.
- `simple_test.py` tests a basic scenario where two clients communicate in a simple channel.
- `utils.py` file that has error messages that you should use.

#### **Important note:**

Your options are to download a Linux virtual machine, or use the instructional servers through SSH, or use the lab machines to test your code.

Side note: If you are using Windows 10, you can use the Linux subsystem for Windows and run your code from there and it may work.

#### **Deliverables**

1. `client.py`
2. `server.py`.

#### **Requirements:**

##### **1. Server Functionality**

The server should accept a single command line argument that's the port that the server should run on.

Unlike your server in part 0, your server in this part of the assignment must allow many clients to be connected and sending messages concurrently. Each client should have an associated name (so that other connected clients know who each message is from) and channel that they're currently subscribed to. When a client first connects, it won't have an associated name and channel. The first message that the server receives from the client should be used as the client's name.

Future messages from the client to the server can take one of two forms. The first type of message is a control message; control messages always begin with "/". There are three different control messages your server should handle from clients:

**/join <channel>** should add the client to the given channel. Clients can only be in one channel at a time, so if the client is already in a channel, this command should remove the client from that channel. When a client is added to a channel, a message should be broadcasted to all existing members of the channel stating that the client has joined. Similarly, if the client left a channel, a message should be broadcasted stating that the client left.

**/create <channel>** should create a new channel with the given name, and add the client to the given channel. As with the /join call, the client should be removed from any existing channels.

**/list** should send a message back to the client with the names of all current channels, separated by newlines.

The second type of message is normal messages to the client's current channel. All messages that are not preceded by a / are considered normal messages. These messages should be broadcasted to all other clients in the channel, preceded by the client's name in brackets (see the example below). Messages should not be sent to client in different channels. If the client is not currently in a channel, the server should send an error message back to the client.

When a client disconnects, a message should be broadcasted to all members of the client's channel saying that the client disconnected.

Sockets provide a datastream functionality, but they don't delineate different messages. When a given `recv` call returns some data, the socket won't tell you whether the data returned is a single message, or multiple messages, or part of one message. As a result, you'll need a way to determine when a message ends. For this assignment, use fixed length messages that have 200 characters for all messages (including messages from the server to the client). If a message is shorter than 200 characters, you should pad the message with spaces (and the receiver should strip any spaces off of the end of the message). You can assume that no messages are longer than 200 characters.

Be sure that your code correctly handles the case where less than one message is available in the socket's buffer (so a `recv` call returns fewer than 200 characters of data) and the case where more than one message is available in the socket's buffer. You should handle partial messages by buffering: if a `recv` call returns only part of a message, your code should hold on to the part of the message until the remainder of the message is received, and then handle the complete message. For example, if a client receives 150 characters from the server, it should hold those 150 characters until 50 more characters are received.

The client should only write the message to stdout once all 200 characters have been received. To help you check for your server's handling of these cases, we've provided a special client (`client_split_messages.py`) that splits messages into many smaller messages before sending them to the server. This client only tests some of the scenarios your server should handle! You'll likely want to modify this client to test for additional cases.

## **2. Error Handling**

Your server should handle cases where the client sends an invalid message by returning an appropriate error message to the client. For example, if a client uses the `/join` command but doesn't provide the name of a channel to join, the server should send back an error message. The provided `utils.py` includes format strings for all of the errors you should handle. You can use these format strings using Python's string formatting operations. For example, `utils.py` defines the following error message:

```
CLIENT_SERVER_DISCONNECTED = "Server at {0}:{1} has disconnected"
```

You can use the `.format` function to replace the brackets with strings as follows:

```
error_message = CLIENT_SERVER_DISCONNECTED.format("localhost", 12345)
```

You are required to use the error messages defined in `utils.py`. If you do not use these error messages (with appropriate formatting), you will not get credit for error handling.

When commands lead to an error, the command should not cause any changes at the server. For example, if a client is currently in the `AA` channel and then tries to join a channel that doesn't exist, the client should not be removed from the `AA` channel.

We will only test for errors that have associated error messages in `utils.py` in our testing. You're welcome to check for additional errors if you'd like, but you will not be graded on them.

## **3. Client Functionality**

Each client connects to a particular chat server, and is associated with a name. Your client should be started as follows:

```
$ python client.py Scott 127.0.0.1 55555
```

This command should connect to the server at the given address and port number, and then send a message with the name `Scott`.

After being started, the client should listen for messages from the server and from stdin. Messages from the server should be printed to the command line (after being stripped of any spaces at the end) and messages from stdin should be sent to the server (after being padded with spaces, as needed). The client should print `[Me]` to each new line, to make it clear which messages in the history were typed by the client. When the client gets a message from the server, it should write over the `[Me]` with the message from the server. For an example of this, take a look at the demo video linked above.

Here's an example of a client's interaction with a server that was started locally on port 55555:

```
python chatv3_client.py Panda localhost 55555
```

```
[Me] Hello world!
```

Not currently in any channel. Must join a channel before sending messages.

```
[Me] /list
```

```
[Me] /create 168_tas
```

```
[Me] /list
```

```
168_tas
```

```
[Me] Hello world!
```

```
Alice has joined
```

```
[Alice] Hi everyone! Does anyone know what we're doing on the first day of lecture?
```

After seeing Alice's message, Panda stopped his client. After Panda created the 168\_tas channel, a second client started:

```
python chatv3_client.py Alice 127.0.0.1 55555
```

```
[Me] /list
```

```
168_tas
```

```
[Me] /join 168_tas
```

```
[Me] Hi everyone! Does anyone know what we're doing on the first day of lecture?
```

```
Panda has left
```

#### **4. Some helpful hints:**

In Python, you can create a socket and connect to a remote endpoint by using the socket library as follows:

```
import socket
```

```
# The socket constructor accepts a few arguments; the defaults are fine for this class.
```

```
client_socket = socket.socket()
```

```
client_socket.connect(("1.2.3.4", 5678))
```

```
client_socket.sendall("Hello World")
```

The example above created a socket and connected it to port 5678 at IP address 1.2.3.4. Then, it sent a "Hello World" message to the server at 1.2.3.4:5678.

The example above created a client socket that was connected to exactly one remote endpoint. When you create a server, you'll typically want to allow multiple remote clients to connect, and you don't usually know the address of those clients when the socket is created. As a result, server sockets work differently:

```
server_socket = socket.socket()
```

```
server_socket.bind(("1.2.3.4", 5678))
```

```
server_socket.listen(5)
```

After creating the socket, rather than connecting to a particular remote destination, the code above bound the socket to a particular IP address and port, which essentially tells the operating system to associate the given IP address and port with the socket. Finally, the listen call listens for connections made to the socket. When a new client connects to the socket, the socket library will create a new socket to use to communicate with that client, so that the server socket can continue to be used to wait for inbound connections from other clients:

```
(new_socket, address) = server_socket.accept()
```

This call blocks until a client connects (using a `connect()` call, as in the example above), and then returns a newly created socket, `new_socket`, that can be used to send and receive data to and from the client. For example, the call

```
message = new_socket.recv(1024)
```

will block until there is data to receive from the client, and will return up to 1024 bytes of data.

Use The Python Socket Programming HOWTO <https://docs.python.org/2/howto/sockets.html> for additional help.